

TITANIC

March 5, 2025

1 Titanic dataset analysis

1.0.1 This homework involves cleaning and show high level analysis of the Titanic dataset. The dataset is available in Kaggle. The dataset contains information about the passengers of the Titanic, such as their age, class, fare, etc.

1.0.2 1. Import and clean the dataset

```
[86]: # Import libraries
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import warnings
warnings.simplefilter("ignore", category=SyntaxWarning)
warnings.simplefilter("ignore", category=FutureWarning)
```

```
[87]: titanic_df = pd.read_csv("../assets/Datos Titanic/datoslimpios.csv",
    encoding="latin1", on_bad_lines="warn")
titanic_df.head()
```

```
[87]: PassengerId  Survived  Pclass    Name \
0             1         0         3    Braund
1             2         1         1  Cumings
2             3         1         3 Heikkinen
3             4         1         1  Futrelle
4             5         0         3     Allen

                                Lastname   Sex  Age  SibSp  Parch  \
0                        Mr. Owen Harris  male  22.0     1     0
1  Mrs. John Bradley (Florence Briggs Thayer) female  38.0     1     0
2                        Miss. Laina     female  26.0     0     0
3  Mrs. Jacques Heath (Lily May Peel)     female  35.0     1     0
4                        Mr. William Henry   male  35.0     0     0

Ticket    Fare Embarked
```

0	A/5 21171	7.2500	S
1	PC 17599	71.2833	C
2	STON/O2. 3101282	7.9250	S
3	113803	53.1000	S
4	373450	8.0500	S

```
[88]: titanic_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Lastname        891 non-null   object
5   Sex             891 non-null   object
6   Age             891 non-null   float64
7   SibSp           891 non-null   int64
8   Parch           891 non-null   int64
9   Ticket          891 non-null   object
10  Fare            891 non-null   float64
11  Embarked        891 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[89]: titanic_df.columns
```

```
[89]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Lastname', 'Sex', 'Age',
        'SibSp', 'Parch', 'Ticket', 'Fare', 'Embarked'],
        dtype='object')
```

```
[90]: numeric_columns = titanic_df.select_dtypes(include=['int64', 'float64']).columns
numeric_df = titanic_df[numeric_columns]
numeric_df.describe()
```

```
[90]:
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	891.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.385152	0.523008	
std	257.353842	0.486592	0.836071	13.259656	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	21.000000	0.000000	
50%	446.000000	0.000000	3.000000	30.000000	0.000000	
75%	668.500000	1.000000	3.000000	35.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
[91]: df_numeric = titanic_df[['Survived', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare']]
```

```
[92]: tendencia_central = df_numeric.describe().applymap(lambda x: f"{x:0.3f}")
tendencia_central
```

```
[92]:
```

	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000	891.000	891.000	891.000	891.000	891.000
mean	0.384	2.309	29.385	0.523	0.382	32.204
std	0.487	0.836	13.260	1.103	0.806	49.693
min	0.000	1.000	0.420	0.000	0.000	0.000
25%	0.000	2.000	21.000	0.000	0.000	7.910
50%	0.000	3.000	30.000	0.000	0.000	14.454
75%	1.000	3.000	35.000	1.000	0.000	31.000
max	1.000	3.000	80.000	8.000	6.000	512.329


```
[93]: df_numeric.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    891 non-null    int64
1   Pclass      891 non-null    int64
2   Age         891 non-null    float64
3   SibSp       891 non-null    int64
4   Parch       891 non-null    int64
5   Fare        891 non-null    float64
dtypes: float64(2), int64(4)
memory usage: 41.9 KB
```

```
[94]: corr_matrix = df_numeric.corr(method='pearson')
# Print corr matrix as a pretty chart of big size

fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(10, 10))
sns.heatmap(corr_matrix, annot=True, cbar=False, annot_kws = {"size":
    ↪8}, vmin=-1, vmax=1, center=0,
cmap=sns.diverging_palette(20, 220, n=200), square=True, ax=ax)
```

```
ax.set_xticklabels(ax.get_xticklabels(),rotation = 45,horizontalalignment = 'right',)
ax.tick_params(labelsize = 10)
```




```
[96]: # Plot frequency distribution of each column in df_numeric
df_numeric['Pclass'] = df_numeric['Pclass'].astype('int')
df_numeric['Pclass'].plot.hist(bins=100)
```

```
/var/folders/cg/l549l5jn7ql8v5_67ml6_p2m0000gn/T/ipykernel_9573/2914623054.py:2:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
`df_numeric['Pclass'] = df_numeric['Pclass'].astype('int')`

[96]: <Axes: ylabel='Frequency'>



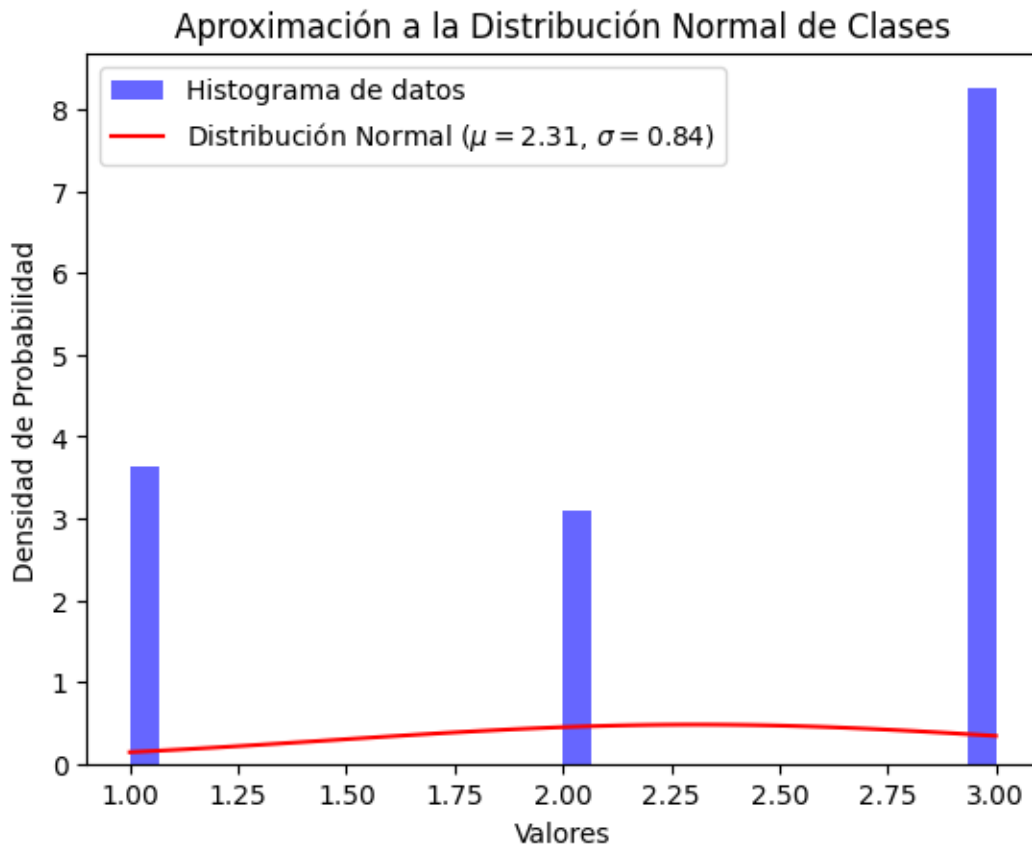
```
[97]: # Estimar parámetros de la distribución normal
mu, sigma = np.mean(df_numeric['Pclass']), np.std(df_numeric['Pclass'])

# Crear el rango de valores para la curva
x = np.linspace(min(df_numeric['Pclass']), max(df_numeric['Pclass']), 100)

y = stats.norm.pdf(x, mu, sigma)

# Graficar el histograma y la curva de densidad
plt.hist(df_numeric['Pclass'], bins=30, density=True, alpha=0.6, color='b',
        label='Histograma de datos')
plt.plot(x, y, 'r', label=f'Distribución Normal ( $\mu={mu:.2f}$ ), ( $\sigma={sigma:.2f}$ )')
plt.xlabel('Valores')
plt.ylabel('Densidad de Probabilidad')
```

```
plt.title('Aproximación a la Distribución Normal de Clases')
plt.legend()
plt.show()
```




```
[98]: lista_valores = df_numeric['Pclass'].to_list()
resultados = {x: lista_valores.count(x) for x in lista_valores}
list(resultados.keys())
```

```
[98]: [3, 1, 2]
```


```
[99]: plt.figure(figsize=(8, 6))
plt.pie(list(resultados.values()), labels=list(resultados.keys()), autopct='%1.
    ↪1f%%', startangle=140,
    ↪colors=['#ff9999', '#66b3ff', '#99ff99', '#ffcc99', '#c2c2f0'])
plt.title("Distribución de Categorías")
plt.show()
```

Distribución de Categorías



```
[100]: df_numeric['Survived'].plot.hist(bins=100)
```

```
[100]: <Axes: ylabel='Frequency'>
```

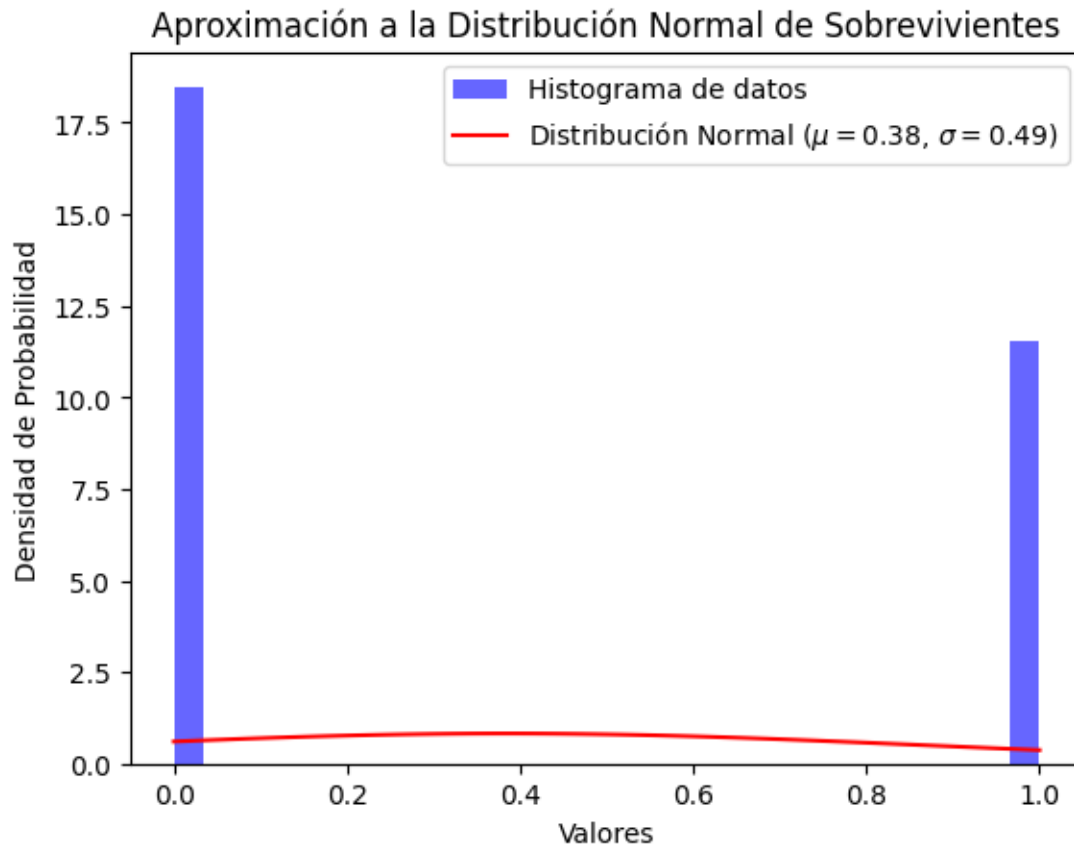


```
[101]: # Estimar parámetros de la distribución normal
mu, sigma = np.mean(df_numeric['Survived']), np.std(df_numeric['Survived'])

# Crear el rango de valores para la curva
x = np.linspace(min(df_numeric['Survived']), max(df_numeric['Survived']), 100)


y = stats.norm.pdf(x, mu, sigma)

# Graficar el histograma y la curva de densidad
plt.hist(df_numeric['Survived'], bins=30, density=True, alpha=0.6, color='b',
        label='Histograma de datos')
plt.plot(x, y, 'r', label=f'Distribución Normal ( $\mu={mu:.2f}$ ), ( $\sigma={sigma:.2f}$ )')
plt.xlabel('Valores')
plt.ylabel('Densidad de Probabilidad')
plt.title('Aproximación a la Distribución Normal de Sobrevivientes')
plt.legend()
plt.show()
```

```
[102]: lista_valores = df_numeric['Survived'].to_list()
resultados = {x: lista_valores.count(x) for x in lista_valores}
list(resultados.keys())
plt.figure(figsize=(8, 6))
plt.pie(list(resultados.values()), labels=list(resultados.keys()), autopct='%1.
    ↪1f%', startangle=140,
    ↪colors=['#ff9999', '#66b3ff', '#99ff99', '#ffcc99', '#c2c2f0'])
plt.title("Distribución de Categorías")
plt.show()
```

Distribución de Categorías



```
[103]: import scipy.stats as stats
```

```
[104]: # Estimar parámetros de la distribución normal


mu, sigma = np.mean(df_numeric['Age']), np.std(df_numeric['Age'])

# Crear el rango de valores para la curva
x = np.linspace(min(df_numeric['Age']), max(df_numeric['Age']), 100)

y = stats.norm.pdf(x, mu, sigma)


# Graficar el histograma y la curva de densidad
plt.hist(df_numeric['Age'], bins=30, density=True, alpha=0.6, color='b',
        label='Histograma de datos')
plt.plot(x, y, 'r', label=f'Distribución Normal ( $\mu={mu:.2f}$ ), ( $\sigma={sigma:.2f}$ )')
plt.xlabel('Valores')
```

```
plt.ylabel('Densidad de Probabilidad')
plt.title('Aproximación a la Distribución Normal de Edad')
plt.legend()
plt.show()
```




```
[105]: df_numeric['Age'].plot.hist(bins=100)
```

```
[105]: <Axes: ylabel='Frequency'>
```



```
[106]: # Plot frequency distribution of each column in df_numeric  
df_numeric['SibSp'].plot.hist(bins=100)
```

```
[106]: <Axes: ylabel='Frequency'>
```




```
[107]: # Estimar parámetros de la distribución normal
mu, sigma = np.mean(df_numeric['SibSp']), np.std(df_numeric['SibSp'])

# Crear el rango de valores para la curva
x = np.linspace(min(df_numeric['SibSp']), max(df_numeric['SibSp']), 100)


y = stats.norm.pdf(x, mu, sigma)

# Graficar el histograma y la curva de densidad
plt.hist(df_numeric['SibSp'], bins=30, density=True, alpha=0.6, color='b',
        label='Histograma de datos')
plt.plot(x, y, 'r', label=f'Distribución Normal ( $\mu={mu:.2f}$ ), ( $\sigma={sigma:.2f}$ )')
plt.xlabel('Valores')
plt.ylabel('Densidad de Probabilidad')
plt.title('Aproximación a la Distribución Normal de Hermanos/Esposos')
plt.legend()
plt.show()
```



```
[108]: # Plot frequency distribution of each column in df_numeric  
df_numeric['Parch'].plot.hist(bins=100)
```

```
[108]: <Axes: ylabel='Frequency'>
```

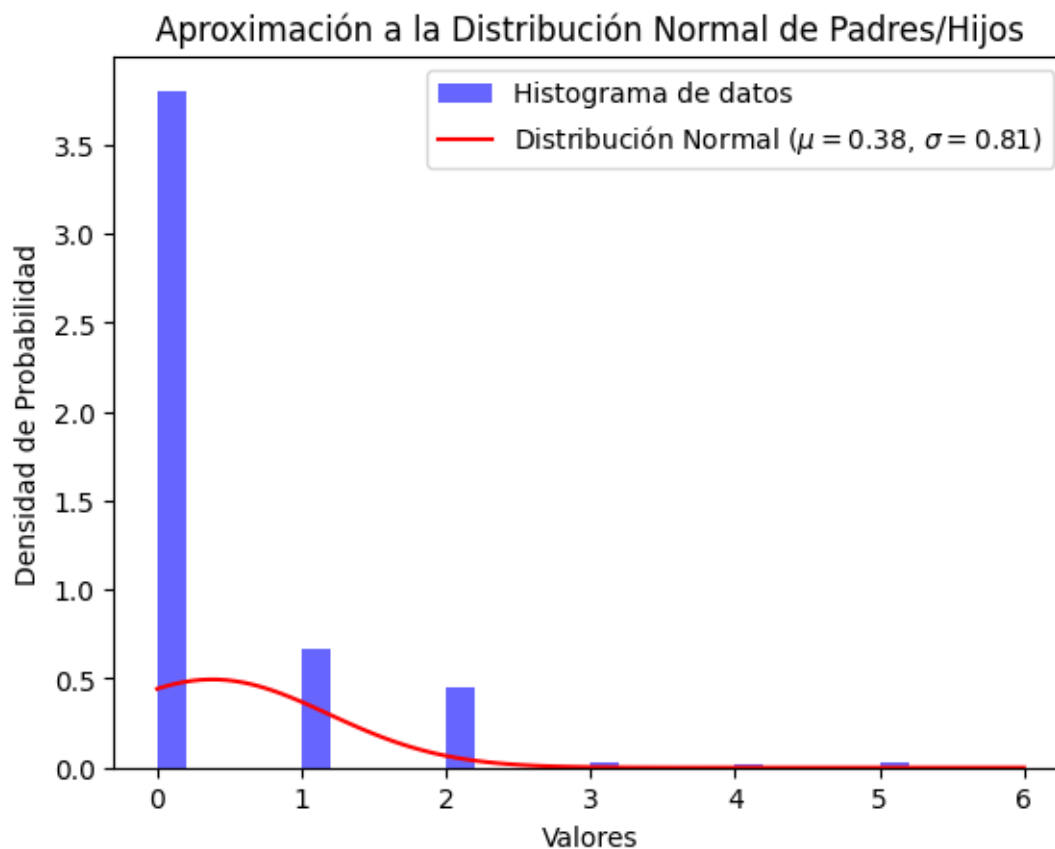


```
[109]: # Estimar parámetros de la distribución normal
mu, sigma = np.mean(df_numeric['Parch']), np.std(df_numeric['Parch'])

# Crear el rango de valores para la curva
x = np.linspace(min(df_numeric['Parch']), max(df_numeric['Parch']), 100)


y = stats.norm.pdf(x, mu, sigma)

# Graficar el histograma y la curva de densidad
plt.hist(df_numeric['Parch'], bins=30, density=True, alpha=0.6, color='b',
        label='Histograma de datos')
plt.plot(x, y, 'r', label=f'Distribución Normal ( $\mu={mu:.2f}$ ), ( $\sigma={sigma:.2f}$ )')
plt.xlabel('Valores')
plt.ylabel('Densidad de Probabilidad')
plt.title('Aproximación a la Distribución Normal de Padres/Hijos')
plt.legend()
plt.show()
```



```
[110]: # Plot frequency distribution of each column in df_numeric  
df_numeric['Fare'].plot.hist(bins=100)
```

```
[110]: <Axes: ylabel='Frequency'>
```





```
[111]: # Estimar parámetros de la distribución normal
mu, sigma = np.mean(df_numeric['Fare']), np.std(df_numeric['Fare'])


# Crear el rango de valores para la curva
x = np.linspace(min(df_numeric['Fare']), max(df_numeric['Fare']), 100)

y = stats.norm.pdf(x, mu, sigma)

# Graficar el histograma y la curva de densidad
plt.hist(df_numeric['Fare'], bins=30, density=True, alpha=0.6, color='b',
        label='Histograma de datos')
plt.plot(x, y, 'r', label=f'Distribución Normal ( $\mu={mu:.2f}$ ), ( $\sigma={sigma:.2f}$ )')
plt.xlabel('Valores')
plt.ylabel('Densidad de Probabilidad')
plt.title('Aproximación a la Distribución Normal en Tarifa')
plt.legend()
plt.show()
```



```
[112]: # Boxplot of first 10 numeric columns
boxplot_df = numeric_df[['Age']]
fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(10, 10))
boxplot_df.boxplot(ax=ax)
ax.set_xticklabels(ax.get_xticklabels(),rotation = 45,horizontalalignment='right',)
ax.tick_params(labelsize = 10)
```



[]: